

Ss

State-space
Search

Heuristic Search

CSE 415: Introduction to Artificial Intelligence
University of Washington
Winter, 2020

Outline

Motivation, Definition

A* Algorithm

Admissibility and Consistency

Heuristics for the 8 Puzzle

Designing Heuristics

Motivation

- Blind search can waste time and space. (due to the combinatorial explosion.)
- Additional knowledge MAY be available -
 - how could it help?
 - E.g., finding a shortest route from Wash. U. to U. Wash. (St. Louis to Seattle).
 - Blind search considers all directions equally, including towards Wash. D.C.
 - Additional knowledge: we need to head northwest.

Ss

State-space
Search

Starting from St. Louis--Blind



Ss

State-space
Search

Starting from St. Louis--Informed



Definition

A *heuristic function* (or simply heuristic) is a function $h: \Sigma \rightarrow \mathbb{R}$, that takes a state as its argument and returns a real number that is an estimate of the distance (or cost) from that state to the closest (or having lowest-cost path) goal state.

$$h(s) = r$$

The function h is typically based on *partial information* about the relationship between each state s and the closest goal state γ to s .

For example, if each state has an (x,y) location, then knowing only x_s and x_γ , we could estimate the distance between s and g as $|x_s - x_\gamma|$.

Outline

Motivation, Definition

A Algorithm*

Admissibility and Consistency

Heuristics for the 8 Puzzle

Designing Heuristics

A* Algorithm

Given a state space Σ having a distance (or cost) function on moves (graph edges): $d(s_i, s_j)$, the A* algorithm searches for a shortest (lowest-cost) path from the initial state s_0 to a goal state γ .

The following algorithm gives the general control structure for A*. It omits a few details:

1. Back pointers for backtracing a path when a goal state is reached.
2. Details of computing g . (done in a manner similar to that in Dijkstra's algorithm, i.e., Uniform Cost Search).
3. Details of implementing the OPEN list and its methods for inserting, finding, and removing.


A* Algorithm

1. For the start state s_0 , compute $f(s_0) = g(s_0) + h(s_0) = h(s_0)$ and put $[s_0, f(s_0)]$ on a list (priority queue) OPEN.
2. If OPEN is empty, output “DONE” and stop.
3. Find and remove the item $[s, p]$ on OPEN **having lowest** p . Break ties arbitrarily.
Put $[s, p]$ on CLOSED.
If s is a goal state: output its description (and backtrace a path), and
if h is known to be admissible, halt.
4. Generate the list L of $[s', f(s')]$ pairs where the s' are the successors of s and their f values are computed using $f(s') = g(s') + h(s')$.
Consider each $[s', f(s')]$.
If there is already a pair $[s', q]$ on CLOSED (for any value q):
if $f(s') > q$, then remove $[s, f(s')]$ from L .
If $f(s') \leq q$, then remove $[s', q]$ from CLOSED.
Else if there is already a pair $[s', q]$ on OPEN (for any value q):
if $f(s') > q$, then remove $[s, f(s')]$ from L .
If $f(s') \leq q$, then remove $[s', q]$ from OPEN.
5. Insert all members of L onto OPEN.
6. Go to Step 2.

A* Algorithm Behavior

During the search A* gives highest priority to that as-yet unexplored state (except in cases where some previously explored state needs to be re-examined) that has the lowest sum of distance from the initial state plus estimated distance to a goal.

Example




state s :	s_0	A	B	C	D	E	F	G	H	γ
heuristic $h(s)$:	14	15	4	10	3	2	16	10	5	0

We show newly enqueued $[s,p]$ pairs in green, and updated $[s,p]$ pairs in red.

Example

OPEN: $[s_0, 14]$




CLOSED:

state s:	s_0	A	B	C	D	E	F	G	H	γ
heuristic $h(s)$:	14	15	4	10	3	2	16	10	5	0

We show newly enqueued $[s,p]$ pairs in green, and updated $[s,p]$ pairs in red.

Example

OPEN: $\{s_0, 14\}$

[B,7], [D,11], [G,14], [A,16], [F,18]

CLOSED: $\{s_0, 14\}$


state s:	s ₀	A	B	C	D	E	F	G	H	γ
heuristic h(s):	14	15	4	10	3	2	16	10	5	0

We show newly enqueued [s,p] pairs in green, and updated [s,p] pairs in red.

Ss

State-space
Search


Example

OPEN: $\{s_0, 14\}$ ~~$\{B, 7\}, \{D, 11\}, \{G, 14\}, \{A, 16\}, \{F, 18\}$~~ $\{D, 9\}, \{G, 14\}, \{A, 16\}, \{F, 18\}, \{C, 22\}$ CLOSED: $\{s_0, 14\}, \{B, 7\}$

state s:	s_0	A	B	C	D	E	F	G	H	γ
heuristic $h(s)$:	14	15	4	10	3	2	16	10	5	0

We show newly enqueued $\{s, p\}$ pairs in green, and updated $\{s, p\}$ pairs in red.

Example



OPEN: $\{s_0, 14\}$


~~$\{B, 7\}, \{D, 11\}, \{G, 14\}, \{A, 16\}, \{F, 18\}$~~
 ~~$\{D, 9\}, \{G, 14\}, \{A, 16\}, \{F, 18\}, \{C, 22\}$~~
 $\{H, 13\}, \{G, 14\}, \{A, 16\}, \{F, 18\}, \{C, 22\}$

CLOSED: $\{s_0, 14\}, \{B, 7\}, \{D, 9\}$

state s:	s_0	A	B	C	D	E	F	G	H	γ
heuristic $h(s)$:	14	15	4	10	5	2	16	8	5	0

We show newly enqueued $\{s, p\}$ pairs in green, and updated $\{s, p\}$ pairs in red.

Example



OPEN: $\{s_0, 14\}$


~~$\{B, 7\}, \{D, 11\}, \{G, 14\}, \{A, 16\}, \{F, 18\}$~~
 ~~$\{D, 9\}, \{G, 14\}, \{A, 16\}, \{F, 18\}, \{C, 22\}$~~
 ~~$\{H, 13\}, \{G, 14\}, \{A, 16\}, \{F, 18\}, \{C, 22\}$~~
 $\{E, 13\}, \{G, 14\}, \{A, 16\}, \{F, 18\}, \{\gamma, 20\},$
 $\{C, 22\}$

CLOSED: $\{s_0, 14\}, \{B, 7\}, \{D, 9\}, \{H, 13\}$

state s:	s ₀	A	B	C	D	E	F	G	H	γ
heuristic h(s):	14	15	4	10	3	2	16	10	5	0

We show newly enqueued $\{s, p\}$ pairs in green, and updated $\{s, p\}$ pairs in red.

Example



OPEN: $[s_0, 14]$

~~$[B, 7], [D, 11], [G, 14], [A, 16], [F, 18]$~~
 ~~$[D, 9], [G, 14], [A, 16], [F, 18], [C, 22]$~~
 ~~$[H, 13], [G, 14], [A, 16], [F, 18], [C, 22]$~~
 ~~$[E, 13], [G, 14], [A, 16], [F, 18], [\gamma, 20],$~~
 ~~$[C, 22]$~~

$[\gamma, 14], [G, 14], [A, 16], [F, 18], [C, 22]$

CLOSED: $[s_0, 14], [B, 7], [D, 9],$
 $[H, 13], [E, 13]$


state s:	s_0	A	B	C	D	E	F	G	H	γ
heuristic $h(s)$:	14	15	4	10	3	2	16	10	5	0

We show newly enqueued $[s, p]$ pairs in green, and updated $[s, p]$ pairs in red.

Ss

State-space
Search

Example

OPEN: $[s_0, 14]$ ~~$[B, 7], [D, 11], [G, 14], [A, 16], [F, 18]$~~ ~~$[D, 9], [G, 14], [A, 16], [F, 18], [C, 22]$~~ ~~$[H, 13], [G, 14], [A, 16], [F, 18], [C, 22]$~~ ~~$[E, 13], [G, 14], [A, 16], [F, 18], [\gamma, 20],$~~ ~~$[C, 22]$~~ ~~$[\gamma, 14], [G, 14], [A, 16], [F, 18], [C, 22]$~~ $[G, 14], [A, 16], [F, 18], [C, 22]$

Predecessor links
that mark the shortest
path

CLOSED: $[s_0, 14], [B, 7], [D, 9],$ $[H, 13], [E, 13], [\gamma, 14]$

state s:	s ₀	A	B	C	D	E	F	G	H	γ
heuristic h(s):	14	15	4	10	3	2	16	10	5	0

We show newly enqueued $[s, p]$ pairs in green, and updated $[s, p]$ pairs in red.

Notes on the Example

The path found is the same as that found by UCS (Dijkstra).

However, fewer nodes are expanded; e.g., A and F are never expanded.

The heuristic h is admissible, so as soon as γ becomes the current state, we can stop.

But h is not consistent: $h(G) - h(D) > d(G, D)$.

This fact doesn't cause trouble here, fortunately.

(If h is consistent, no state will ever have to be expanded a second time, i.e., never have to be moved from CLOSED back to OPEN.)

Outline

Motivation, Definition

A* Algorithm

Admissibility and Consistency

Heuristics for the 8 Puzzle

Designing Heuristics

Admissibility and Consistency

Heuristic h is admissible if and only if

$$(\forall s \in \Sigma)(\forall \gamma \in \Gamma) h(s) \leq d(s, \gamma)$$

Here $d(s, \gamma)$ is the length (cost) of the shortest (lowest-cost) path from s to γ .

I.e., h estimates but never overestimates the distance from s to the closest goal.

If h is admissible, then A^* , using h , finds a shortest path (optimal path) as soon as it expands any goal γ .

Admissibility and Consistency

Heuristic h is *consistent* if and only if

For each edge (s_i, s_j) in the problem-space graph,
$$h(s_i) - h(s_j) \leq d(s_i, s_j)$$

Here $d(s_i, s_j)$ is the length (cost) of the edge from s_i to s_j .

If h is consistent, then along any shortest path from a node (state) s to its closest goal, then h values will be monotonically non-increasing along the path.

If h is consistent, then A*, using h , never has to re-expand a node.

Admissibility of A^*

A^* is *admissible*: Provided its heuristic is admissible, and a path exists from s_0 to γ , then A^* will find a shortest path and stop.

Outline

Motivation, Definition

A* Algorithm

Admissibility and Consistency

Heuristics for the 8 Puzzle

Designing Heuristics

Heuristics for the Eight Puzzle

$h_0(s) = 0$ (uninformed; blind search)

$h_1(s)$ = number of tiles out of place. ("Hamming") $0 \leq h_1(s) \leq 8$.

$\gamma =$

	1	2
3	4	5
6	7	8

$s_a =$

3	1	2
7	6	5
4		8

$h_1(s_a) = 4$.

Heuristics for the Eight Puzzle

$h_2(s)$ = number of rows tile 7 is away from its place.

$\gamma =$

	1	2
3	4	5
6	7	8

$s_a =$

3	1	2
7	6	5
4		8

$$h_2(s_a) = 1.$$

Using tile 7 is arbitrary here -- just an example of a heuristic based on a single, particular tile.

Heuristics for the Eight Puzzle

$h_3(s)$ = number of rows **and** **columns** tile 7 is away from its place.

 $\gamma =$

	1	2
3	4	5
6	7	8

 $s_a =$

3	1	2
7	6	5
4		8

$$h_3(s_a) = 2.$$

This sum is known as the *Manhattan distance* (for a single tile).

Heuristics for the Eight Puzzle

$h_4(s)$ = sum of Manhattan distances for all 8 tiles.

$\gamma =$

	1	2
3	4	5
6	7	8

$s_a =$

3	1	2
7	6	5
4		8

$$h_4(s_a) = 7.$$

This is called the Manhattan distance heuristic.

In this example $h_4(s_a) = h(s_a)$ (the actual shortest distance).

Heuristics for the Eight Puzzle

$$h_5(s) = \text{sum of Euclidean distances for all 8 tiles.}$$

$$= \sum_{i=1}^8 \sqrt{dx_i^2 + dy_i^2}$$

 $\gamma =$

	1	2
3	4	5
6	7	8

 $s_a =$

3	1	2
7	6	5
4		8

$$h_5(s_a) = 1 + 3\sqrt{2} \approx 5.2326$$

This is called the Euclidean distance heuristic.

In (at least) this example $h_5(s_a) < h_4(s_a)$.

Euclidean is not as good as Manhattan.

Heuristic Domination

If $(\forall s \in \Sigma) h_i(s) \geq h_j(s)$, then we say h_i dominates $h_j(s)$.
However, we assume both heuristics are admissible.

If h_i dominates h_j , then we call h_i "more informed" than h_j .
Having a highly informed heuristic is good for limiting a search to relevant parts of the state space.

However, one has to trade off this off against the higher computational cost that usually goes with more informed heuristics.

Outline

Motivation, Definition

A* Algorithm

Admissibility and Consistency

Heuristics for the 8 Puzzle

Designing Heuristics

Designing Heuristics

A common approach to defining heuristics is to create a simpler type of problem.

E.g., For the 8 Puzzle, allow tiles to be removed and put back anywhere, and "charge" a cost of 0.5 for each removal and 0.5 for each putting back. That leads to the Hamming heuristic.

Or: allow tiles to be piled up on top of one another, thus making it easier to move each tile (still one square at a time) to its destination. This leads to the Manhattan heuristic.

Designing Heuristics (cont)

Another way to simplify: change some of the tiles into "blanks" (like the blank tile in Scrabble). The new goal is to get only the non-blank tiles into their proper positions; the blanks are "don't-care" tiles that still take up space, but whose relative ordering is not important. For example:

$$\gamma' =$$

	1	2
3	4	5
/	/	/

$$s_a' =$$

3	1	2
/	/	5
4		/

To compute $h(s_a')$, we transform s_a into s_a' , and solve the simplified problem, getting a path length $d(s_a', \gamma')$, which we use as the value of h . If the reduced problem is easy enough, then we can precompute a table of $d(s_a', \gamma')$ values to speed up computing h during the search. Such a table is called a *pattern database*.

Outline

Motivation, Definition

A* Algorithm

Admissibility and Consistency

Heuristics for the 8 Puzzle

Completeness and Optimality

Designing Heuristics

done!